

# Software Systems

Day 4 – Hex, Memory, and Pointers

# Agenda

- Announcements
- Hex
- Memory
- Pointers

# Announcements

- Lots of people seem to have issues building code.
- Much of this is likely due to being on Ubuntu 20 instead of Ubuntu 22.
- You can update in-place, but read up on the subtleties of it. (Things can still break because package names can change.)
- I'm investigating possible fixes.
  - Please post in Discord for help, and also include as much information about your setup as you can.
  - I'll try to fix issues for the release of Assignment 1.

# Today: Hex, Memory, Pointers

- Systems is about understanding what a computer is doing under the hood.
- Compared to Python, C gives you a much closer view of what the machine is doing.
- There are still abstractions, but it's a lower-level view.

# Today: Hex, Memory, Pointers

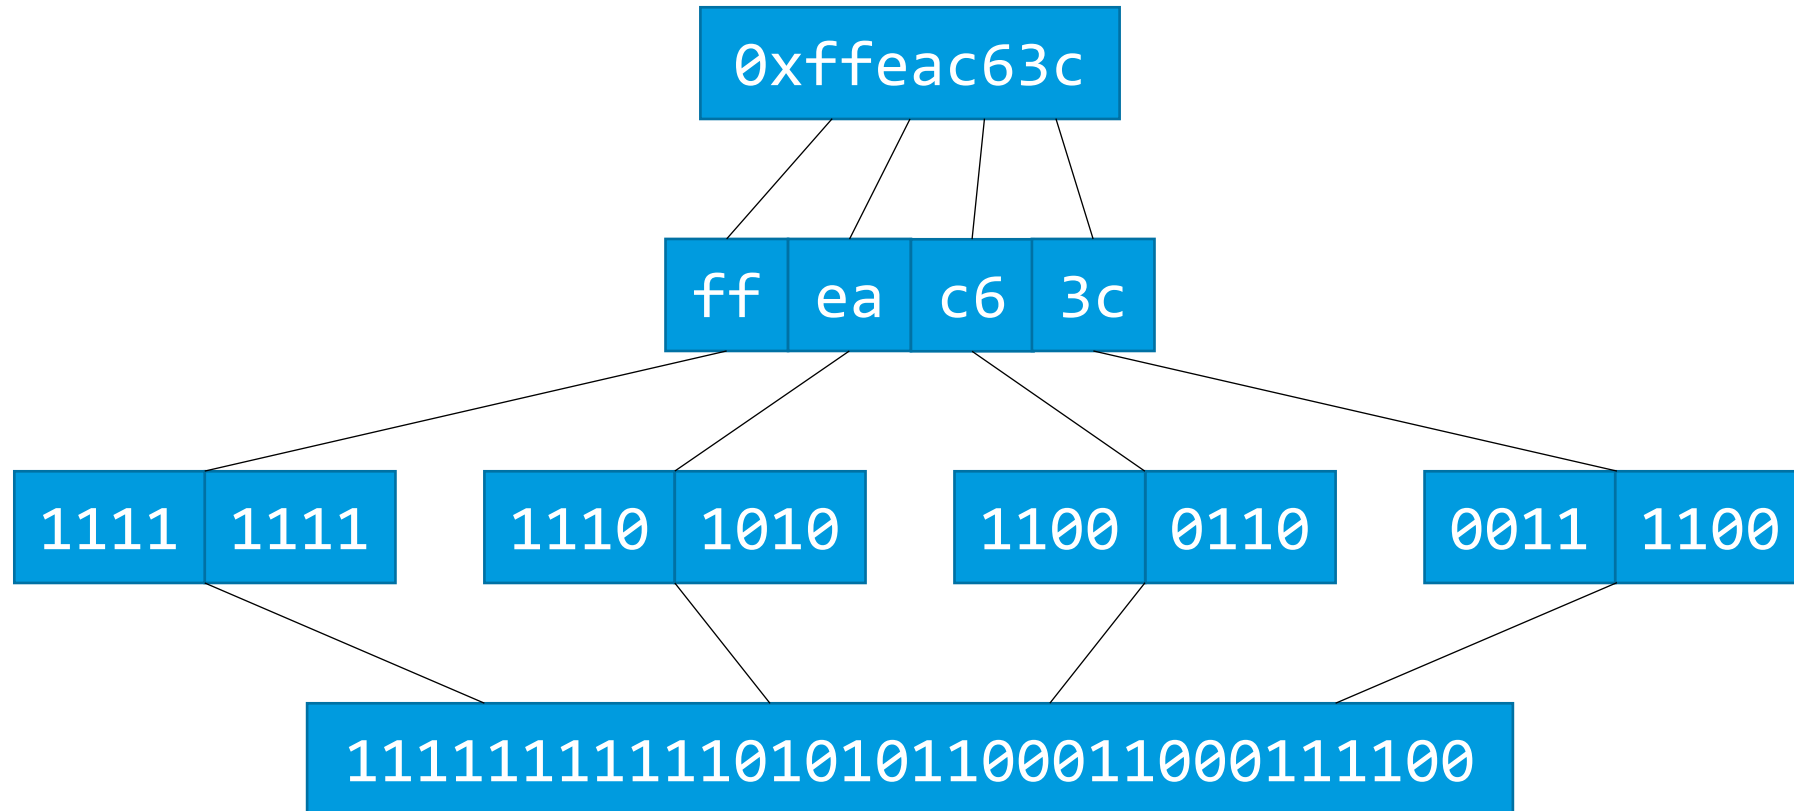
- With compilation, we saw that
  - Variable names don't get compiled into the program.
  - Things are instead identified by where in memory they live.
- Today, we'll learn:
  - Hexadecimal notation and why it is useful.
  - The memory layout and how it's used in a program.
  - Pointers and how they help us work with memory in C.

# Hexadecimal Notation (Hex)

- Everything that a computer works with is ultimately in bits.
- Bits are written as binary digits - 1s and 0s.
  - The number 42 is 101010.
  - The string "sus" in C is 01110011011101010111001100000000.
- But this is tedious to read.
- Hexadecimal works in powers of 16 rather than powers of 2.
  - Group 4 bits at a time.
  - Still small enough to work with, but much more compact.

# Hexadecimal Notation (Hex)

- 32-bit:  $2^{32}$  addresses, 64-bit:  $2^{64}$  addresses.



# Hexadecimal Notation (Hex)

- Hex is often written with the prefix `0x` to distinguish it from other notations.
- Usually written as an even number of digits, with leading 0s if necessary (e.g., `0x0f` instead of `0xf`).



# Hexadecimal Notation (Hex)

- Let's do an exercise with hex. A hexercise, if you will.
- What is 0xb in binary? In decimal?
- What is 0xbeef in decimal?
- What is 0xcafe in binary?
- What is the string "IYKYK" in binary? In hex?
  - Look up an ASCII table for how to encode characters.
- What is the longest word you can spell with only hexadecimal digits?

# Memory

- Head First C Chapter 2 talks about the memory layout.
- From top to bottom:
  - Stack (functions and local variables)
  - Heap (dynamically allocated memory)
  - Globals (variables accessible to everyone)
  - Constants (values that can't be changed)
  - Code (the compiled/linked bytes of the program)

# Memory

- In reality, this isn't quite accurate - the kernel takes part of memory.
- Specifically, the upper half.
- But actually, this isn't quite accurate either.
- 64 bits of address space is a lot of memory.
  - About 16 million TB can be referenced with 64 bits.
- These are not the real addresses in memory.
  - There's a translation step (for next time).



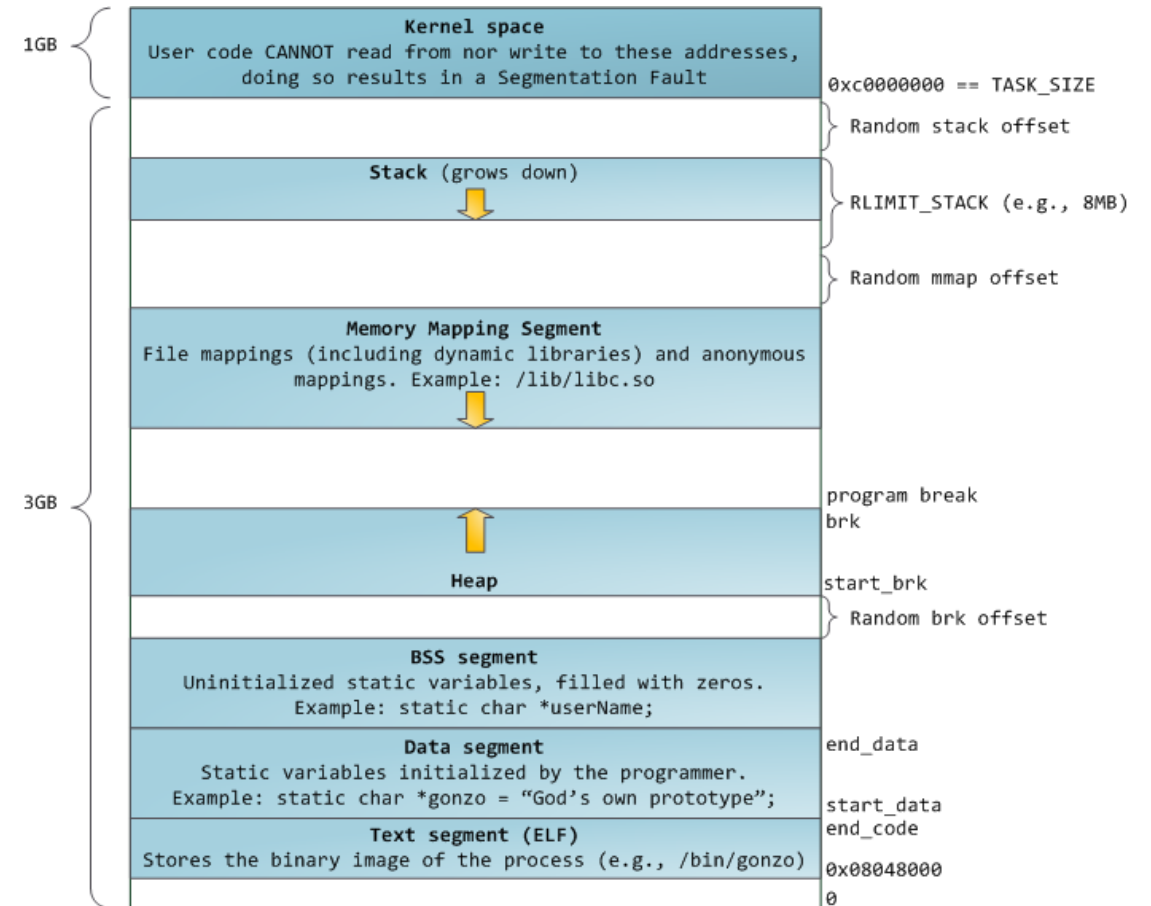
# Memory

- To save on memory translation costs, the upper 17 bits must be identical for all "canonical" (valid) addresses in memory.
- So only the lower 48 bits are used as addresses.
- Exercise:
  - What is the lowest "higher half" address?
  - What is the highest "lower half" address?
  - How much total memory can be expressed with 48 bits?



# Memory

- In practice, the layout is a bit more complicated than what Head First C shows.
- The stack grows down, and the heap grows up.



# Pointers

- Pointers are how we refer to, and access, locations in memory.
- Declare pointers using \*:
  - `int* p;`
- Dereference pointers also using \*:
  - `int x = *p;`
- Find the location of a variable using &:
  - `int* p = &x;`

# Pointers

- You can also combine dereference and assignment:
  - `*p = 42;`
- And `a[3]` is just "syntactic sugar" for `*(a + 3)`.
- This means that `3[a]` is also completely valid (though cursed).

# Pointers

- To print a pointer, use %p:  

```
int x = 42;  
printf("%p\n", &x);
```
- Exercise: try allocating pointers in different parts of memory, and printing them.
  - To allocate memory on the heap:  

```
#include <stdlib.h>  
// Do stuff here  
int* p = malloc(sizeof(int));  
// Do stuff here  
free(p);
```
  - What ranges of addresses do you see?



# Pointers and Types

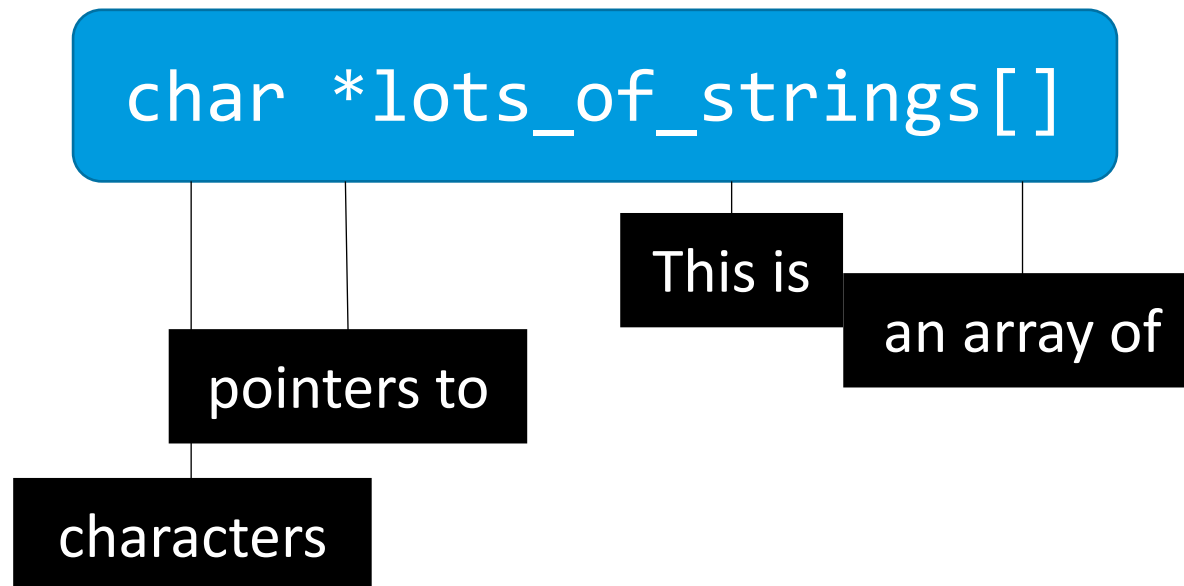
- Why does C care about what kind of data it points to?
- Dereferencing a pointer "removes" the \*:  

```
int x = 42;  
int* p = &x;  
int** pp = &p;  
int* y = *pp;
```
- C needs to know what it can do with the data after dereferencing.
- For arrays, the pointer arithmetic needs to work out:  

```
int numbers[3] = {3, 1, 4}; // Usually 12 bytes  
int a = numbers[1]; // 4 bytes from numbers
```
- You can also have a generic pointer: `void*` (but don't use it now).

# Pointers and Types

- So how do you know what type `char *lots_of_strings[]` is?
- Use the right-to-left rule:
  - Start from the variable name and read right.
  - Then go back to the variable name and read left.



# Pointers and Types

- Exercise: suppose that you have:  
`int* x = malloc(sizeof(int));`  
`int a[3] = {3, 1, 4};`  
`const char* name = "Steve";`
- What type is each expression?
  - `&x` - `int**`
  - `a[1]` - `int`
  - `*x` - `int`
  - `*name` - `(const char)`
  - `name + 1` - `const char*`

# Pointers: Passing to Functions

- Everything in C is what we call pass by value:
  - If you pass something to a function, it makes a copy of the value.
- What does the following program print?

```
#include <stdio.h>
```

```
void foo(int x) {  
    x = 42;  
}
```

```
int main(void) {  
    int x = 0;  
    foo(x);  
    printf("%d\n", x);  
    return 0;  
}
```

# Pointers: Passing to Functions

- If you want to let other functions change a variable, you can provide a pointer.
- What does the following program print?

```
#include <stdio.h>
```

```
void foo(int* p) {  
    *p = 42;  
}
```

```
int main(void) {  
    int x = 0;  
    foo(&x);  
    printf("%d\n", x);  
    return 0;  
}
```

# Pointers: Passing to Functions

- Be aware that Python always passes by pointer under the hood.

- So this will print [1, 3, 4]:

```
def sort(L):  
    L.sort()
```

```
my_list = [3, 1, 4]  
sort(my_list)
```

# Pointers

- <https://www.youtube.com/watch?v=mnXkiAKbUPg>
- <https://www.youtube.com/watch?v=G7LJC9vJluU>